

# Introduction to dplyr

*Jenny Bryan, modified by Roger H. French*

*29 September, 2018*

## Contents

5.2.2.1	Intro . . . . .	1
5.2.2.1.1	Load <b>dplyr</b> . . . . .	1
5.2.2.1.2	Load the Gapminder data . . . . .	2
5.2.2.2	Meet <b>tbl_df</b> , an upgrade to <b>data.frame</b> . . . . .	2
5.2.2.3	Think before you create excerpts of your data . . . . .	3
5.2.2.4	Use <b>filter()</b> to subset data row-wise. . . . .	4
5.2.2.5	Meet the new pipe operator . . . . .	5
5.2.2.6	Use <b>select()</b> to subset the data on variables or columns. . . . .	6
5.2.2.7	Revel in the convenience . . . . .	7
5.2.2.8	Pause to reflect . . . . .	8
5.2.2.9	Links . . . . .	9

### 5.2.2.1 Intro

**dplyr** is a new package for data manipulation.

- It is part of the tidyverse,
  - and is loaded with the tidyverse metapackage
- It is built to be fast, highly expressive, and open-minded
  - about how your data is stored.
- It is developed by Hadley Wickham and Romain Francois.

**dplyr**'s roots are in an earlier, still-very-useful package

- called **plyr**,
- which implements the “split-apply-combine” strategy for data analysis.

Where **plyr** covers a diverse set of inputs and outputs

- (e.g., arrays, data.frames, lists),
- **dplyr** has a laser-like focus on data.frames and related structures.

Have no idea what I'm talking about?

- Not sure if you care?
- If you use these base R functions:
  - **subset()**, **apply()**, **[sl]apply()**, **tapply()**,
  - **aggregate()**, **split()**, **do.call()**,
- then you should keep reading.

#### 5.2.2.1.1 Load dplyr

```
## install if you do not already have
```

```
## from CRAN:
```

```
## install.packages('dplyr')
```

```
## from GitHub using devtools (which you also might need to install!):
```

```
## devtools::install_github("hadley/lazyeval")
## devtools::install_github("hadley/dplyr")
## suppressPackageStartupMessages(library(dplyr))
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

### 5.2.2.1.2 Load the Gapminder data

An excerpt of the Gapminder data which we work with alot.

```
gd_url <- "http://www.stat.ubc.ca/~jenny/notOcto/STAT545A/examples/gapminder/data/gapminderDataFiveYear"
# gd_url <- "http://tiny.cc/gapminder"
gdf <- read.delim(file = gd_url)
str(gdf)
```

```
## 'data.frame':   1704 obs. of  6 variables:
## $ country   : Factor w/ 142 levels "Afghanistan",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ year      : int   1952 1957 1962 1967 1972 1977 1982 1987 1992 1997 ...
## $ pop       : num   8425333 9240934 10267083 11537966 13079460 ...
## $ continent: Factor w/  5 levels "Africa","Americas",...: 3 3 3 3 3 3 3 3 3 3 ...
## $ lifeExp   : num   28.8 30.3 32 34 36.1 ...
## $ gdpPercap: num   779 821 853 836 740 ...
```

```
head(gdf)
```

```
##      country year      pop continent lifeExp gdpPercap
## 1 Afghanistan 1952  8425333      Asia  28.801   779.4453
## 2 Afghanistan 1957  9240934      Asia  30.332   820.8530
## 3 Afghanistan 1962 10267083      Asia  31.997   853.1007
## 4 Afghanistan 1967 11537966      Asia  34.020   836.1971
## 5 Afghanistan 1972 13079460      Asia  36.088   739.9811
## 6 Afghanistan 1977 14880372      Asia  38.438   786.1134
```

### 5.2.2.2 Meet tibble, an upgrade to data.frame

```
gtbl <- tibble(gdf)
gtbl
```

```
## # A tibble: 1,704 x 6
##   country      year      pop continent lifeExp gdpPercap
##   <fct>      <int>    <dbl> <fct>      <dbl>    <dbl>
## 1 Afghanistan 1952  8425333 Asia      28.8      779.
## 2 Afghanistan 1957  9240934 Asia      30.3      821.
## 3 Afghanistan 1962 10267083 Asia      32.0      853.
## 4 Afghanistan 1967 11537966 Asia      34.0      836.
```

```
## 5 Afghanistan 1972 13079460 Asia 36.1 740.
## 6 Afghanistan 1977 14880372 Asia 38.4 786.
## 7 Afghanistan 1982 12881816 Asia 39.9 978.
## 8 Afghanistan 1987 13867957 Asia 40.8 852.
## 9 Afghanistan 1992 16317921 Asia 41.7 649.
## 10 Afghanistan 1997 22227415 Asia 41.8 635.
## # ... with 1,694 more rows
```

```
glimpse(gtbl)
```

```
## Observations: 1,704
## Variables: 6
## $ country <fct> Afghanistan, Afghanistan, Afghanistan, Afghanistan, ...
## $ year <int> 1952, 1957, 1962, 1967, 1972, 1977, 1982, 1987, 1992...
## $ pop <dbl> 8425333, 9240934, 10267083, 11537966, 13079460, 1488...
## $ continent <fct> Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia...
## $ lifeExp <dbl> 28.801, 30.332, 31.997, 34.020, 36.088, 38.438, 39.8...
## $ gdpPercap <dbl> 779.4453, 820.8530, 853.1007, 836.1971, 739.9811, 78...
```

A `tbl_df` is basically an improved `data.frame`,

- or a tibble dataframe
- for which `dplyr` provides nice methods for high-level inspection.

Specifically, these methods do something sensible

- for datasets with many observations and/or variables.
- You do **NOT** need to turn your `data.frames`
  - into `tbl_dfs` to use `plyr`.
- I do so here for demonstration purposes only.

### 5.2.2.3 Think before you create excerpts of your data ...

If you feel the urge to store a little snippet of your data:

```
(snippet <- subset(gdf, country == "Canada"))
```

```
##   country year      pop continent lifeExp gdpPercap
## 241 Canada 1952 14785584 Americas 68.750 11367.16
## 242 Canada 1957 17010154 Americas 69.960 12489.95
## 243 Canada 1962 18985849 Americas 71.300 13462.49
## 244 Canada 1967 20819767 Americas 72.130 16076.59
## 245 Canada 1972 22284500 Americas 72.880 18970.57
## 246 Canada 1977 23796400 Americas 74.210 22090.88
## 247 Canada 1982 25201900 Americas 75.760 22898.79
## 248 Canada 1987 26549700 Americas 76.860 26626.52
## 249 Canada 1992 28523502 Americas 77.950 26342.88
## 250 Canada 1997 30305843 Americas 78.610 28954.93
## 251 Canada 2002 31902268 Americas 79.770 33328.97
## 252 Canada 2007 33390141 Americas 80.653 36319.24
```

Stop and ask yourself ...

Do I want to create mini datasets for each level of some factor (or unique combination of several factors) ... in order to compute or graph something?

If YES, use proper data aggregation techniques

- or facetting in `ggplot2` plots

- or conditioning in `lattice`

– **don't subset the data.**

Or, more realistic,

- only subset the data as a temporary measure
- while you develop your elegant code
  - for computing on or visualizing these data subsets.

If NO, then maybe you really do need to store a copy of a subset of the data.

- But seriously consider whether you can achieve your goals
- by simply using the `subset` = argument of,
  - e.g., the `lm()` function,
  - to limit computation to your excerpt of choice.
- Lots of functions offer a `subset` = argument!

Copies and excerpts of your data

- clutter your workspace,
- invite mistakes,
- and sow general confusion.
- Avoid whenever possible.

Reality can also lie somewhere in between.

- You will find the workflows presented below
  - can help you accomplish your goals
  - with minimal creation of temporary, intermediate objects.

#### 5.2.2.4 Use `filter()` to subset data row-wise.

`filter()` takes logical expressions

- and returns the rows for which all are TRUE.

```
filter(gtbl, lifeExp < 29)
```

```
## # A tibble: 2 x 6
##   country    year    pop continent lifeExp gdpPercap
##   <fct>      <int>   <dbl> <fct>      <dbl>    <dbl>
## 1 Afghanistan 1952 8425333 Asia      28.8     779.
## 2 Rwanda      1992 7290203 Africa    23.6     737.
```

```
filter(gtbl, country == "Rwanda")
```

```
## # A tibble: 12 x 6
##   country    year    pop continent lifeExp gdpPercap
##   <fct>      <int>   <dbl> <fct>      <dbl>    <dbl>
## 1 Rwanda    1952 2534927 Africa     40     493.
## 2 Rwanda    1957 2822082 Africa    41.5    540.
## 3 Rwanda    1962 3051242 Africa     43    597.
## 4 Rwanda    1967 3451079 Africa    44.1    511.
## 5 Rwanda    1972 3992121 Africa    44.6    591.
## 6 Rwanda    1977 4657072 Africa     45    670.
## 7 Rwanda    1982 5507565 Africa    46.2    882.
## 8 Rwanda    1987 6349365 Africa    44.0    848.
## 9 Rwanda    1992 7290203 Africa    23.6    737.
## 10 Rwanda   1997 7212583 Africa    36.1    590.
```

```
## 11 Rwanda    2002 7852401 Africa    43.4    786.
## 12 Rwanda    2007 8860588 Africa    46.2    863.
```

```
filter(gtbl, country %in% c("Rwanda", "Afghanistan"))
```

```
## # A tibble: 24 x 6
##   country      year      pop continent lifeExp gdpPercap
##   <fct>      <int>    <dbl> <fct>      <dbl>    <dbl>
## 1 Afghanistan 1952  8425333 Asia      28.8     779.
## 2 Afghanistan 1957  9240934 Asia      30.3     821.
## 3 Afghanistan 1962 10267083 Asia      32.0     853.
## 4 Afghanistan 1967 11537966 Asia      34.0     836.
## 5 Afghanistan 1972 13079460 Asia      36.1     740.
## 6 Afghanistan 1977 14880372 Asia      38.4     786.
## 7 Afghanistan 1982 12881816 Asia      39.9     978.
## 8 Afghanistan 1987 13867957 Asia      40.8     852.
## 9 Afghanistan 1992 16317921 Asia      41.7     649.
## 10 Afghanistan 1997 22227415 Asia      41.8     635.
## # ... with 14 more rows
```

Compare with some base R code

- to accomplish the same things

```
gdf[gdf$lifeExp < 29, ] ## repeat `gdf`, [i, j] indexing is distracting
subset(gdf, country == "Rwanda") ## almost same as filter ... but wait ...
```

### 5.2.2.5 Meet the new pipe operator

Before we go any further,

- we should exploit the new pipe operator
- that `dplyr` imports from the [magrittr](#) package.

This changes your data analytical life.

You no longer need to enact multi-operation commands

- by nesting them inside each other,
- like so many [Russian nesting dolls](#).

This new syntax leads to code

- that is much easier to write and to read.

Here's what it looks like: `%>%`.

The RStudio keyboard shortcut:

- Ctrl + Shift + M (Linux/Windows),
- Cmd + Shift + M (Mac),
- according to [this tweet](#).

Let's demo then I'll explain:

```
gdf %>% head
```

```
##   country year      pop continent lifeExp gdpPercap
## 1 Afghanistan 1952  8425333      Asia  28.801   779.4453
## 2 Afghanistan 1957  9240934      Asia  30.332   820.8530
## 3 Afghanistan 1962 10267083      Asia  31.997   853.1007
```

```
## 4 Afghanistan 1967 11537966      Asia 34.020 836.1971
## 5 Afghanistan 1972 13079460      Asia 36.088 739.9811
## 6 Afghanistan 1977 14880372      Asia 38.438 786.1134
```

This is equivalent to `head(gdf)`.

- This pipe operator takes the thing on the left-hand-side
- and **pipes** it into the function call on the right-hand-side -literally, it drops it in as the first argument.

Never fear, you can still specify other arguments to this function!

To see the first 3 rows of Gapminder,

- we could say `head(gdf, 3)`
- or this:

```
gdf %>% head(3)
```

```
##      country year      pop continent lifeExp gdpPercap
## 1 Afghanistan 1952 8425333      Asia 28.801 779.4453
## 2 Afghanistan 1957 9240934      Asia 30.332 820.8530
## 3 Afghanistan 1962 10267083      Asia 31.997 853.1007
```

I've advised you to think "gets"

- whenever you see the assignment operator, `<-`.

Similarly, you should think "then"

- whenever you see the pipe operator, `%>%`.

You are probably not impressed yet,

- but the magic will soon happen.

### 5.2.2.6 Use `select()` to subset the data on variables or columns.

Back to `dplyr` ...

Use `select()` to subset the data

- on variables or columns.

Here's a conventional call:

```
select(gtobl, year, lifeExp) ## tbl_df prevents TMI from printing
```

```
## # A tibble: 1,704 x 2
##   year lifeExp
##   <int>   <dbl>
## 1 1952    28.8
## 2 1957    30.3
## 3 1962    32.0
## 4 1967    34.0
## 5 1972    36.1
## 6 1977    38.4
## 7 1982    39.9
## 8 1987    40.8
## 9 1992    41.7
## 10 1997    41.8
## # ... with 1,694 more rows
```

And here's similar operation,

- but written with the pipe operator
- and piped through `head`:

```
gtbl %>%
  select(year, lifeExp) %>%
  head(4)
```

```
## # A tibble: 4 x 2
##   year lifeExp
##   <int>   <dbl>
## 1  1952    28.8
## 2  1957    30.3
## 3  1962    32.0
## 4  1967    34.0
```

Think:

- “Take `gtbl`,
- then select the variables `year` and `lifeExp`,
- then show the first 4 rows.”

### 5.2.2.7 Revel in the convenience

Here’s the data for Cambodia,

- but only certain variables:

```
gtbl %>%
  filter(country == "Cambodia") %>%
  select(year, lifeExp)
```

```
## # A tibble: 12 x 2
##   year lifeExp
##   <int>   <dbl>
## 1  1952    39.4
## 2  1957    41.4
## 3  1962    43.4
## 4  1967    45.4
## 5  1972    40.3
## 6  1977    31.2
## 7  1982    51.0
## 8  1987    53.9
## 9  1992    55.8
## 10 1997    56.5
## 11 2002    56.8
## 12 2007    59.7
```

and what a typical base R call would look like:

```
gdf[gdf$country == "Cambodia", c("year", "lifeExp")]
```

```
##   year lifeExp
## 217 1952  39.417
## 218 1957  41.366
## 219 1962  43.415
## 220 1967  45.415
## 221 1972  40.317
## 222 1977  31.220
```

```
## 223 1982 50.957
## 224 1987 53.914
## 225 1992 55.803
## 226 1997 56.534
## 227 2002 56.752
## 228 2007 59.723
```

or, possibly?, a nicer look using base R's `subset()` function:

```
subset(gdf, country == "Cambodia", select = c(year, lifeExp))
```

```
##      year lifeExp
## 217 1952 39.417
## 218 1957 41.366
## 219 1962 43.415
## 220 1967 45.415
## 221 1972 40.317
## 222 1977 31.220
## 223 1982 50.957
## 224 1987 53.914
## 225 1992 55.803
## 226 1997 56.534
## 227 2002 56.752
## 228 2007 59.723
```

#### 5.2.2.8 Pause to reflect

We've barely scratched the surface of `dplyr`

- but I want to point out key principles you may start to appreciate.

`dplyr`'s verbs,

- such as `filter()` and `select()`,
- are what's called [pure functions](#). To quote from Wickham's [Advanced R Programming book](#):

" The functions that are the easiest to understand and reason about

- are pure functions:
  - functions that always map the same input to the same output
  - and have no other impact on the workspace. In other words, pure functions have no side effects:
  - they don't affect the state of the world in any way
  - apart from the value they return."

In fact, these verbs are a special case of pure functions:

- they take the same flavor of object as input and output.
- Namely, a `data.frame` or one of the other data receptacles `dplyr` supports.
- And finally, the data is **always** the very first argument
  - of the verb functions.

This set of deliberate design choices,

- together with the new pipe operator,
- produces a highly effective,
- low friction [domain-specific language](#)
- for data analysis.



#### 5.2.2.9 Links

[Jenny Bryan Stat 545](#)

dplyr official stuff

- package home [on CRAN](#)
  - note there are several vignettes, with the [introduction](#) being the most relevant right now
  - the [one on window functions](#) will also be interesting to you now
- development home [on GitHub](#)
- [tutorial HW delivered](#) (note this links to a DropBox folder) at useR! 2014 conference

Blog post [Hands-on dplyr tutorial for faster data manipulation in R](#) by Data School, that includes a link to an R Markdown document and links to videos

[Cheatsheet](#) I made for `dplyr` join functions (not relevant yet but soon)